

```
In [ ]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
```

```
In [ ]: df = pd.read_csv("heart.csv")
pd.concat([df.head(10), df.tail(10)])
```

	0	40	M	ATA	140	289	0	Normal	172	N	0.0	Up	0
	1	49	F	NAP	160	180	0	Normal	156	N	1.0	Flat	1
	2	37	M	ATA	130	283	0	ST	98	N	0.0	Up	0
	3	48	F	ASY	138	214	0	Normal	108	Y	1.5	Flat	1
	4	54	M	NAP	150	195	0	Normal	122	N	0.0	Up	0
	5	39	M	NAP	120	339	0	Normal	170	N	0.0	Up	0
	6	45	F	ATA	130	237	0	Normal	170	N	0.0	Up	0
	7	54	M	ATA	110	208	0	Normal	142	N	0.0	Up	0
	8	37	M	ASY	140	207	0	Normal	130	Y	1.5	Flat	1
	9	48	F	ATA	120	284	0	Normal	120	N	0.0	Up	0
	908	63	M	ASY	140	187	0	LVH	144	Y	4.0	Up	1
	909	63	F	ASY	124	197	0	Normal	136	Y	0.0	Flat	1
	910	41	M	ATA	120	157	0	Normal	182	N	0.0	Up	0
	911	59	M	ASY	164	176	1	LVH	90	N	1.0	Flat	1
	912	57	F	ASY	140	241	0	Normal	123	Y	0.2	Flat	1
	913	45	M	TA	110	264	0	Normal	132	N	1.2	Flat	1
	914	68	M	ASY	144	193	1	Normal	141	N	3.4	Flat	1
	915	57	M	ASY	130	131	0	Normal	115	Y	1.2	Flat	1
	916	57	F	ATA	130	236	0	LVH	174	N	0.0	Flat	1
	917	38	M	NAP	138	175	0	Normal	173	N	0.0	Up	0

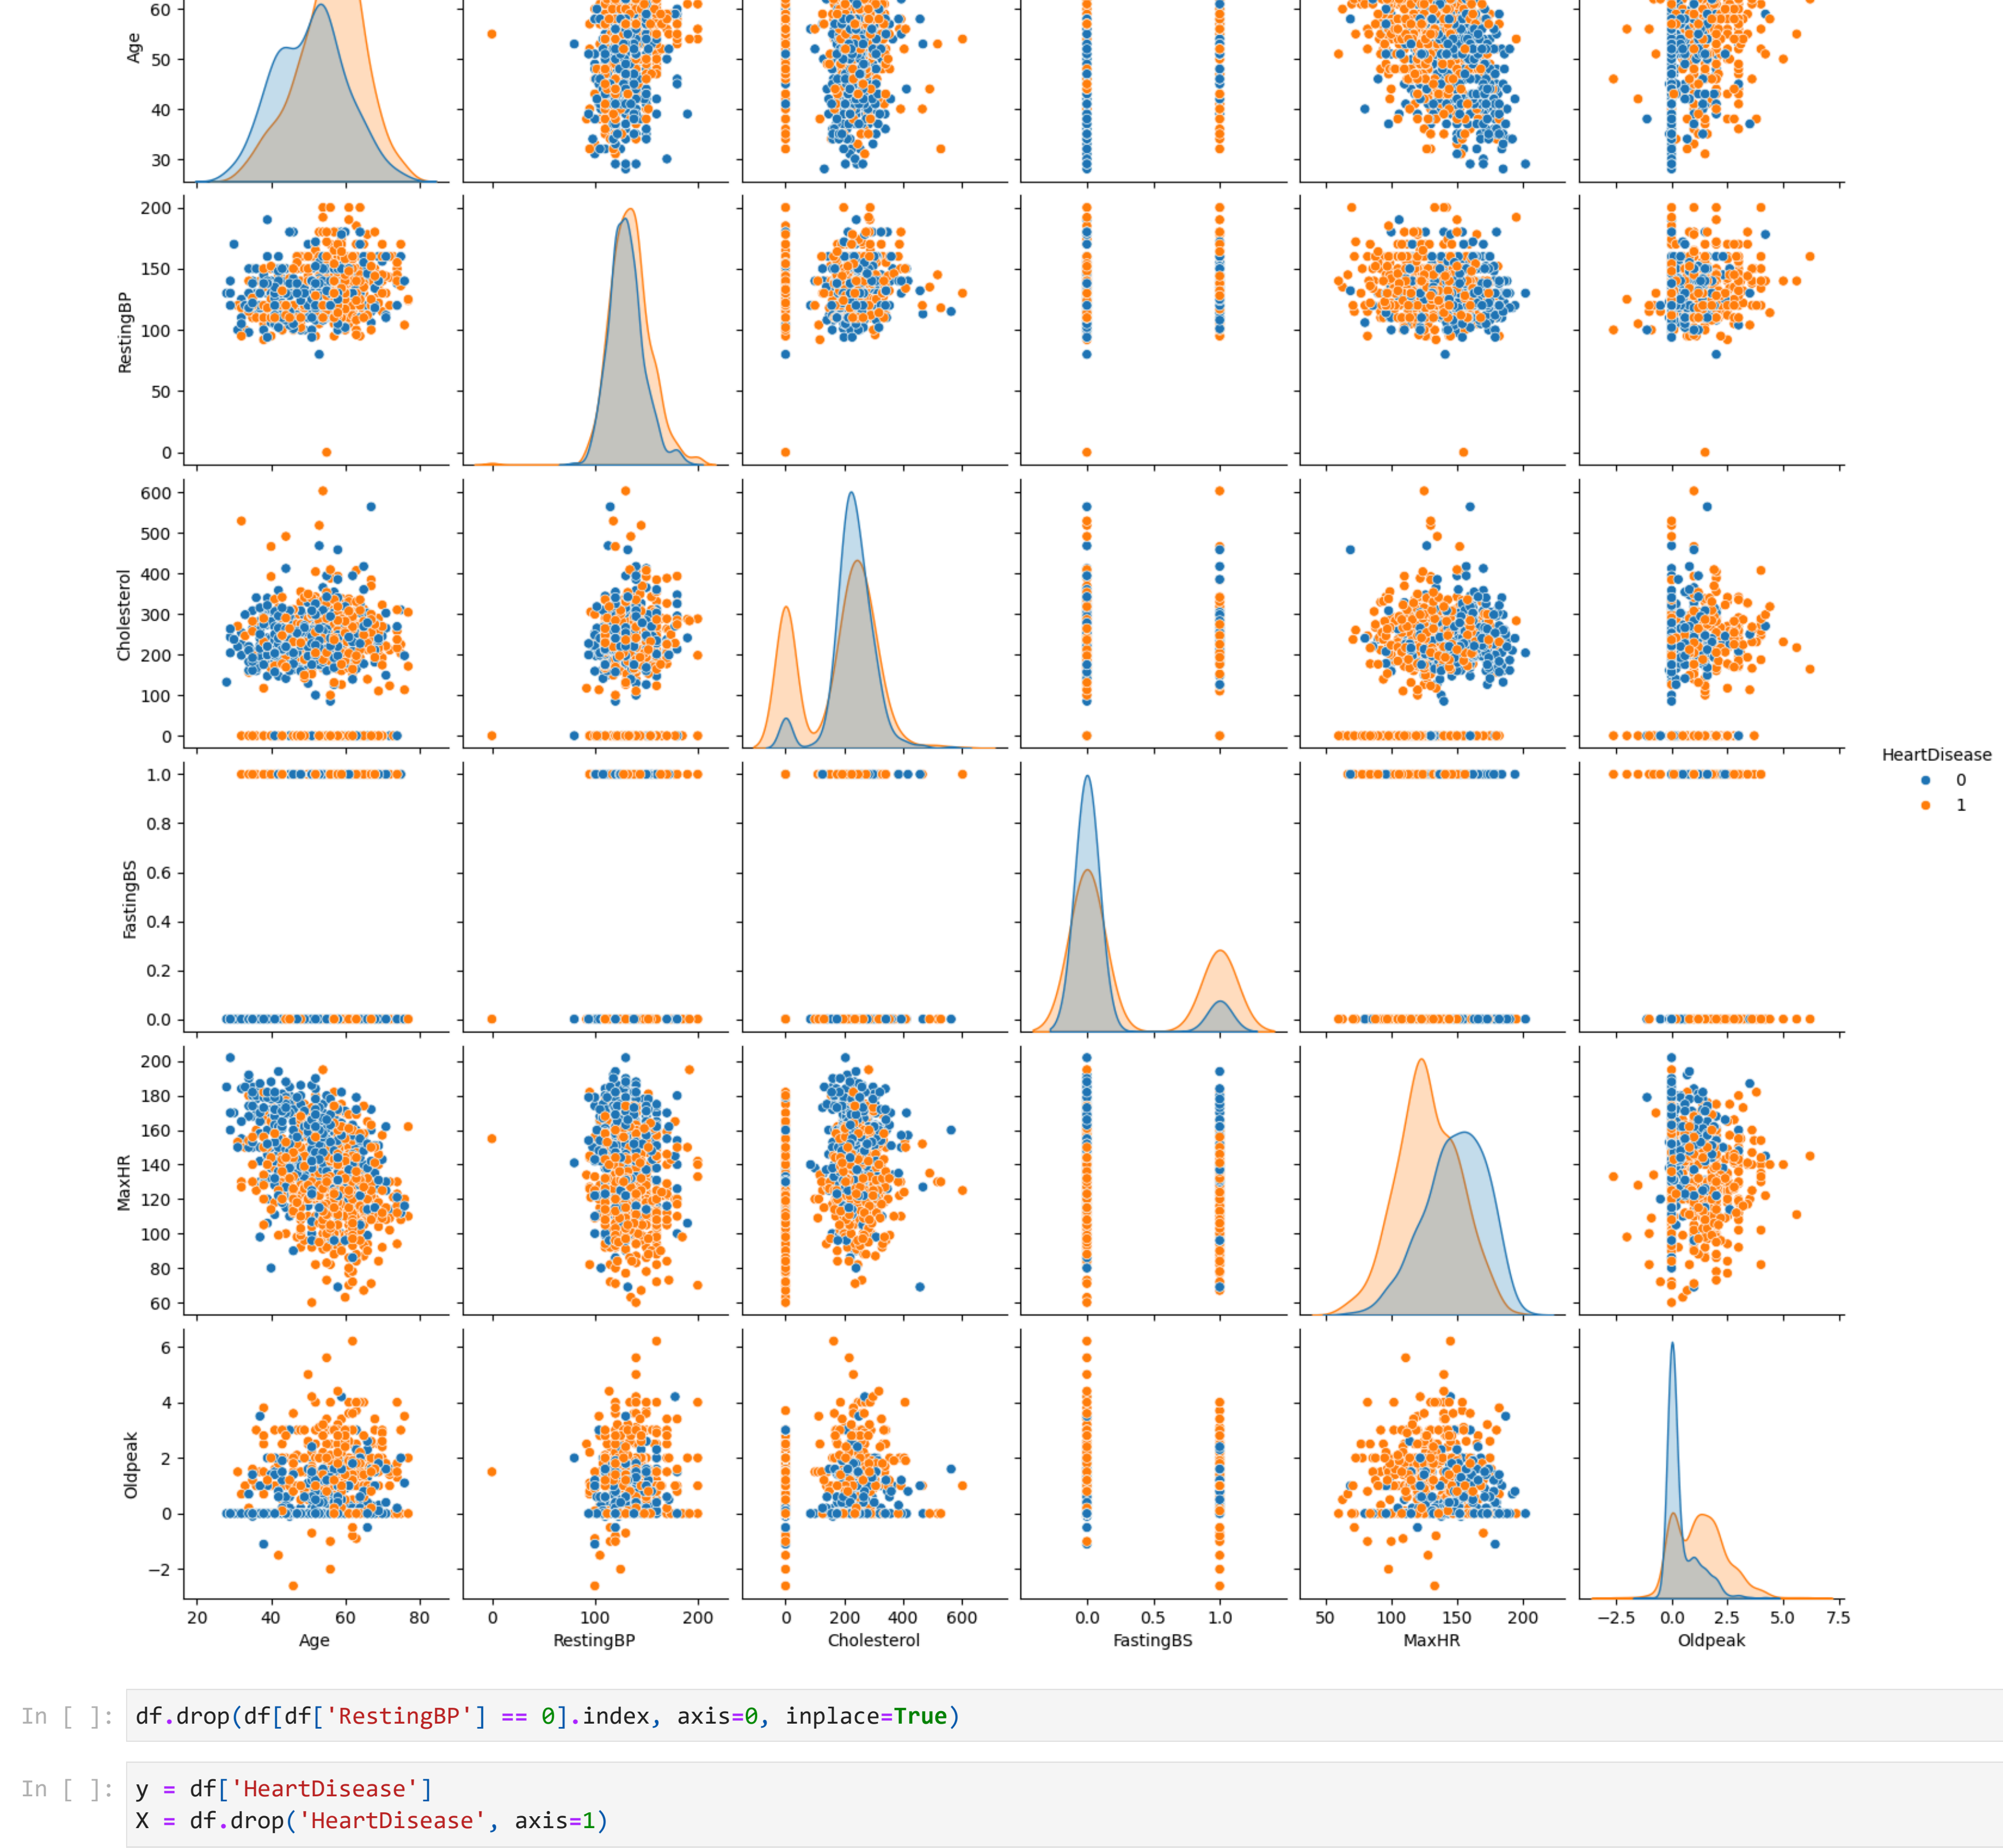
In []: df.describe(include="all")

```
In [ ]: df.describe(include="all")
```

	unique	NaN	2		4	NaN	NaN	NaN	3	NaN		2	NaN	3	
	top	NaN	M		ASY	NaN	NaN	NaN	Normal	NaN		N	NaN	Flat	
	freq	NaN	725		496	NaN	NaN	NaN	552	NaN		547	NaN	460	
	mean	53.510893	NaN		NaN	132.396514	198.799564	0.233115	NaN	136.809368		NaN	0.887364	NaN	0.553
	std	9.432617	NaN		NaN	18.514154	109.384145	0.423046	NaN	25.460334		NaN	1.066570	NaN	0.497
	min	28.000000	NaN		NaN	0.000000	0.000000	0.000000	NaN	60.000000		NaN	-2.600000	NaN	0.000
	25%	47.000000	NaN		NaN	120.000000	173.250000	0.000000	NaN	120.000000		NaN	0.000000	NaN	0.000
	50%	54.000000	NaN		NaN	130.000000	223.000000	0.000000	NaN	138.000000		NaN	0.600000	NaN	1.000
	75%	60.000000	NaN		NaN	140.000000	267.000000	0.000000	NaN	156.000000		NaN	1.500000	NaN	1.000
	max	77.000000	NaN		NaN	200.000000	603.000000	1.000000	NaN	202.000000		NaN	6.200000	NaN	1.000

```
In [ ]: import seaborn as sns
```

```
sns.pairplot(df, hue='HeartDisease')
```



```
In [ ]: df.drop(df[df['RestingBP'] == 0].index, axis=0, inplace=True)
```

```
In [ ]: y = df['HeartDisease']
X = df.drop('HeartDisease', axis=1)
```

```
In [ ]: from sklearn.preprocessing import OneHotEncoder

categoricals = X.select_dtypes('object').columns
integer = X.select_dtypes(['int', 'float']).columns

encoder = OneHotEncoder(handle_unknown='ignore')

X_encoded = encoder.fit_transform(X[categoricals])
X_encoded = pd.DataFrame.sparse.from_spmatrix(X_encoded, columns=encoder.get_feature_names_out(), index=X.index)

X = pd.concat([X[integer], X_encoded], axis=1)
pd.concat([X.head(10), X.tail(10)])
```

Out[]:	Age	RestingBP	Cholesterol	FastingBS	MaxHR	Oldpeak	Sex_F	Sex_M	ChestPainType_ASY	ChestPainType_ATA	ChestPainType_NAP	ChestPainType_TA
	0	40	140	289	0	172	0.0	0	1.0	0	1.0	0
	1	49	160	180	0	156	1.0	1.0	0	0	0	1.0
	2	37	130	283	0	98	0.0	0	1.0	0	1.0	0
	3	48	138	214	0	108	1.5	1.0	0	1.0	0	0
	4	54	150	195	0	122	0.0	0	1.0	0	0	1.0
	5	39	120	339	0	170	0.0	0	1.0	0	0	1.0
	6	45	130	237	0	170	0.0	1.0	0	0	1.0	0
	7	54	110	208	0	142	0.0	0	1.0	0	1.0	0
	8	37	140	207	0	130	1.5	0	1.0	1.0	0	0
	9	48	120	284	0	120	0.0	1.0	0	0	1.0	0
	908	63	140	187	0	144	4.0	0	1.0	1.0	0	0
	909	63	124	197	0	136	0.0	1.0	0	1.0	0	0
	910	41	120	157	0	182	0.0	0	1.0	0	1.0	0
	911	59	164	176	1	90	1.0	0	1.0	1.0	0	0
	912	57	140	241	0	123	0.2	1.0	0	1.0	0	0
	913	45	110	264	0	132	1.2	0	1.0	0	0	0
	914	68	144	193	1	141	3.4	0	1.0	1.0	0	0
	915	57	130	131	0	115	1.2	0	1.0	1.0	0	0
	916	57	130	236	0	174	0.0	1.0	0	0	1.0	0
	917	38	138	175	0	173	0.0	0	1.0	0	0	1.0

```
In [ ]: from sklearn.model_selection import train_test_split
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
X_train.shape, X_test.shape, y_train.shape, y_test.shape
```

```
Out[ ]: ((733, 20), (184, 20), (733,), (184,))
```

```
In [ ]: from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import classification_report

clft = DecisionTreeClassifier(random_state=42)
clft.fit(X_train, y_train)
report = classification_report(y_test, clft.predict(X_test))
for line in str(report).split('\n'):
    print(line)
```

	precision	recall	f1-score	support
0	0.70	0.86	0.77	72
1	0.89	0.76	0.82	112
accuracy			0.80	184
macro avg	0.80	0.81	0.80	184
weighted avg	0.82	0.80	0.80	184

```
x:\_Netology\DS_ModelsParametersModule1\env\Lib\site-packages\sklearn\utils\validation.py:877: UserWarning: pandas.DataFrame with sparse columns found.It will be converted to a dense numpy array.
warnings.warn(
```

```
x:\_Netology\DS_ModelsParametersModule1\env\Lib\site-packages\sklearn\utils\validation.py:877: UserWarning: pandas.DataFrame with sparse columns found.It will be converted to a dense numpy array.
warnings.warn(
```

```
In [ ]: from sklearn.ensemble import RandomForestClassifier
```

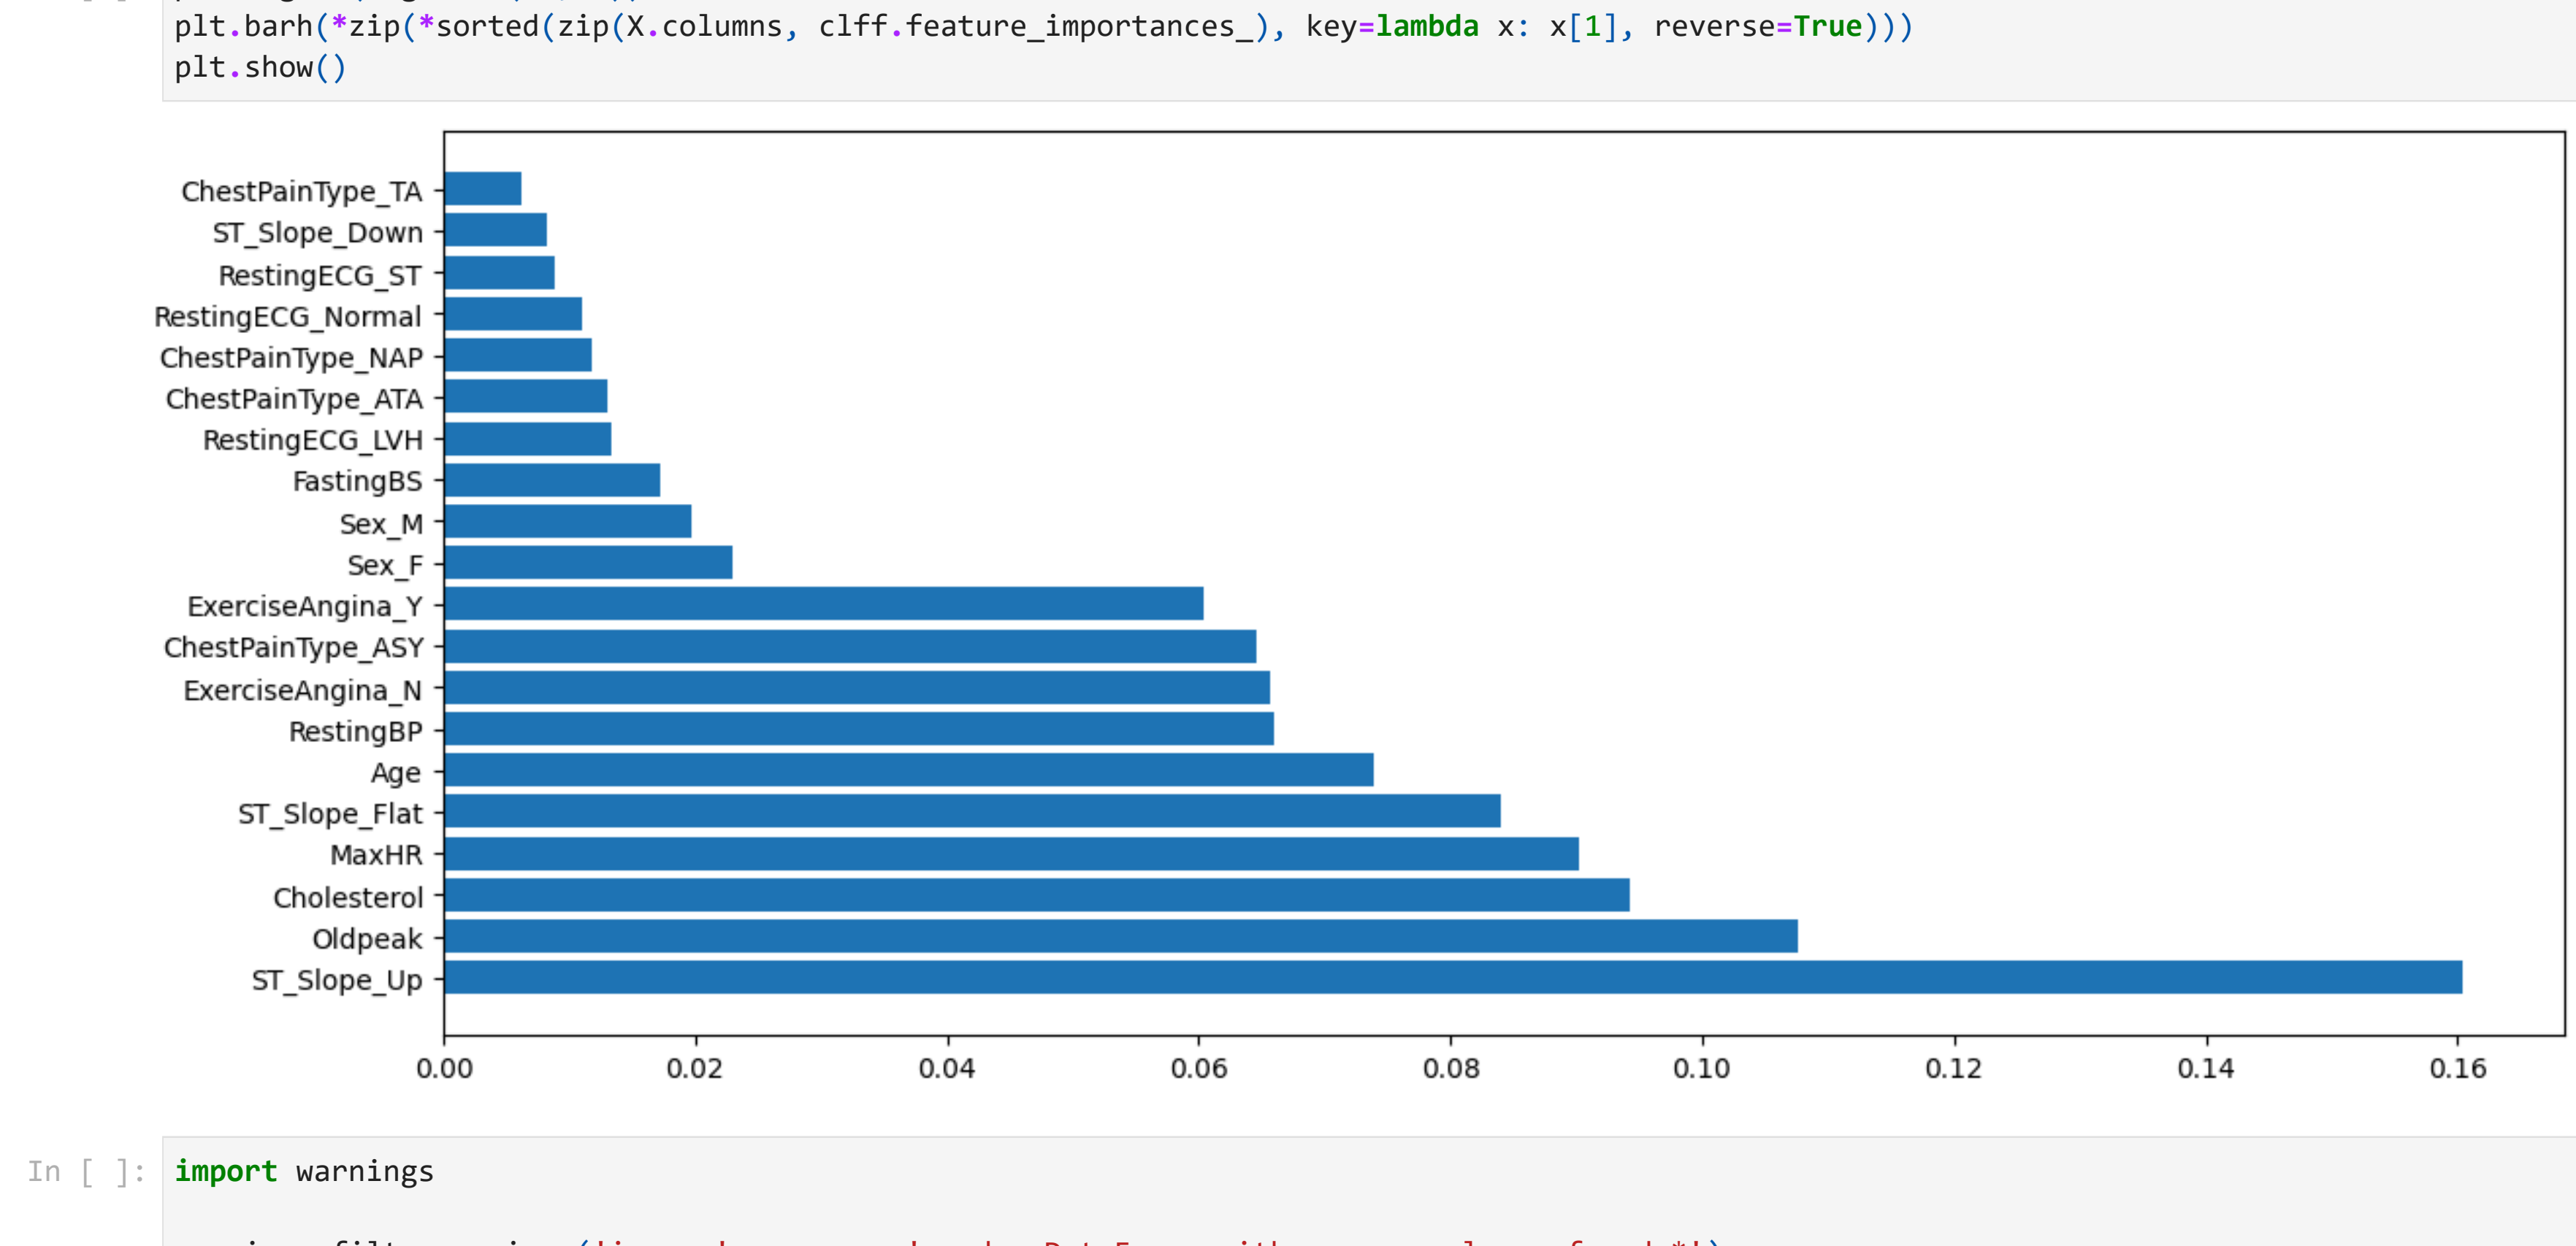
```
clff = RandomForestClassifier(random_state=42)
clff.fit(X_train, y_train)
report = classification_report(y_test, clff.predict(X_test))
for line in str(report).split('\n'):
    print(line)
```

```
x:\_Netology\DS_ModelsParametersModule1\env\Lib\site-packages\sklearn\utils\validation.py:877: UserWarning: pandas.DataFrame with sparse columns found.It will be converted to a dense numpy array.
warnings.warn(
```

	precision	recall	f1-score	support
0	0.85	0.86	0.86	72
1	0.91	0.90	0.91	112
accuracy			0.89	184
macro avg	0.88	0.88	0.88	184
weighted avg	0.89	0.89	0.89	184

```
x:\_Netology\DS_ModelsParametersModule1\env\Lib\site-packages\sklearn\utils\validation.py:877: UserWarning: pandas.DataFrame with sparse columns found.It will be converted to a dense numpy array.
warnings.warn(
```

```
In [ ]: plt.figure(figsize=(14, 6))
plt.barh(*zip(*sorted(zip(X.columns, clff.feature_importances_), key=lambda x: x[1], reverse=True)))
plt.show()
```



```
In [ ]: import warnings
```

```
warnings.filterwarnings('ignore', message='pandas.DataFrame with sparse columns found.**')
```

```
In [ ]: from sklearn.ensemble import BaggingClassifier
```

```
clfb = BaggingClassifier(clft, random_state=42)
clfb.fit(X_train, y_train)
report = classification_report(y_test, clfb.predict(X_test))
for line in str(report).split('\n'):
    print(line)
```

	precision	recall	f1-score	support
0	0.83	0.89	0.86	72
1	0.93	0.88	0.90	112
accuracy			0.89	184
macro avg	0.88	0.89	0.88	184
weighted avg	0.89	0.89	0.89	184

```
In [ ]: from sklearn.ensemble import StackingClassifier
```

```
from sklearn.svm import LinearSVC
from sklearn.linear_model import LogisticRegression
```

```
clfs = StackingClassifier([('DecisionTreeClassifier', clft),
                           ('RandomForestClassifier', clff),
                           ('LinearSVC', LinearSVC(random_state=42))],
                          final_estimator=LogisticRegression(random_state=42))
```

```
clfs.fit(X_train, y_train)
report = classification_report(y_test, clfs.predict(X_test))
for line in str(report).split('\n'):
    print(line)
```

	precision	recall	f1-score	support
0	0.83	0.88	0.85	72
1	0.92	0.88	0.90	112
accuracy			0.88	184
macro avg	0.87	0.88	0.88	184
weighted avg	0.88	0.88	0.88	184

```
In [ ]: from sklearn.ensemble import GradientBoostingClassifier
```

```
clfg = GradientBoostingClassifier(random_state=42)
clfg.fit(X_train, y_train)
report = classification_report(y_test, clfg.predict(X_test))
for line in str(report).split('\n'):
    print(line)
```

	precision	recall	f1-score	support
0	0.86	0.82	0.84	72
1	0.89	0.91	0.90	112
accuracy			0.88	184
macro avg	0.87	0.87	0.87	184
weighted avg	0.87	0.88	0.87	184

```
In [ ]: from catboost import CatBoostClassifier
```

```
clfc = CatBoostClassifier(random_state=42, verbose=0)
clfc.fit(X_train, y_train)
report = classification_report(y_test, clfc.predict(X_test))
for line in str(report).split('\n'):
    print(line)
```

	precision	recall	f1-score	support
0	0.84	0.86	0.85	72
1	0.91	0.89	0.90	112
accuracy			0.88	184
macro avg	0.87	0.88	0.88	184
weighted avg	0.88	0.88	0.88	184